

# Searching in non-deterministic, partially observable and unknown environments

CE417: Introduction to Artificial Intelligence  
Sharif University of Technology  
Spring 2016

Soleymani

# Problem types

---

- ▶ **Deterministic and fully observable** (single-state problem)
  - ▶ Agent knows exactly its state even after a sequence of actions
  - ▶ Solution is a sequence
- ▶ **Non-observable or sensor-less** (conformant problem)
  - ▶ Agent's percepts provide no information at all
  - ▶ Solution is a sequence
- ▶ **Nondeterministic and/or partially observable** (contingency problem)
  - ▶ Percepts provide new information about current state
  - ▶ Solution can be a contingency plan (tree or strategy) and not a sequence
  - ▶ Often interleave search and execution
- ▶ **Unknown state space** (exploration problem)

# More complex than single-state problem

---

- ▶ Searching with nondeterministic actions
- ▶ Searching with partial observations
- ▶ Online search & unknown environment

# Non-deterministic or partially observable env.

---

- ▶ Perception become useful
  - ▶ Partially observable
    - ▶ To narrow down the set of possible states for the agent
  - ▶ Non-deterministic
    - ▶ To show which outcome of the action has occurred
- ▶ Future percepts can not be determined in advance
- ▶ Solution is a contingency plan
  - ▶ A tree composed of nested if-then-else statements
  - ▶ What to do depending on what percepts are received
- ▶ Now, we focus on an agent design that finds a guaranteed plan before execution (not online search)

# Searching with non-deterministic actions

---

- ▶ In **non-deterministic** environments, the result of an action can vary.
  - ▶ Future percepts can specify which outcome has occurred.
- ▶ Generalizing the **transition function**
  - ▶  $RESULTS: S \times A \rightarrow 2^S$  instead of  $RESULTS: S \times A \rightarrow S$
- ▶ Search tree will be an AND-OR tree.
  - ▶ Solution will be a sub-tree containing a contingency plan (nested **if-then-else** statements)

# Erratic vacuum world

- ▶ States


- ▶  $\{1, 2, \dots, 8\}$

- ▶ Actions

- ▶ {Left, Right, Suck}

- ▶ Goal

- ▶  $\{7\}$  or  $\{8\}$




- ▶ Non-deterministic:

- ▶ When sucking a dirty square, it cleans it and sometimes cleans up dirt in an adjacent square.
  - ▶ When sucking a clean square, it sometimes deposits dirt on the carpet.

# AND-OR search tree

**AND node:** environment's choice of outcome

**OR node:** agent's choices of actions



[Suck, if State=5 then [Right, Suck] else []]

# Solution to AND-OR search tree

---

- ▶ **Solution** for AND-OR search problem is a **sub-tree** that:
  - ▶ specifies **one action** at each **OR** node
  - ▶ includes **every outcome** at each **AND** node
  - ▶ **has a goal node at every leaf**
- ▶ Algorithms for searching AND-OR graphs
  - ▶ Depth first
  - ▶ BFS, best first,  $A^*$ , ...



**function** AND-OR-GRAPH-SEARCH(*problem*) **returns** a conditional plan or failure  
OR-SEARCH(*problem*.INITIAL-STATE, *problem*, [ ])

**function** OR-SEARCH(*state*, *problem*, *path*) **returns** a conditional plan or failure  
**if** *problem*.GOAL-TEST(*state*) **then return** the empty plan  
**if** *state* is on *path* **then return** failure  
**for each** *action* **in** *problem*.ACTIONS(*state*) **do**  
    *plan*  $\leftarrow$  AND-SEARCH(RESULTS(*state*, *action*), *problem*, [*state* | *path*])  
    **if** *plan*  $\neq$  failure **then return** [*action* | *plan*]  
**return** failure

**function** AND-SEARCH(*states*, *problem*, *path*) **returns** a conditional plan, or failure

**for each**  $s_i$  in *states* **do**  
    *plan*<sub>*i*</sub>  $\leftarrow$  OR-SEARCH( $s_i$ , *problem*, *path*)  
    **if** *plan*<sub>*i*</sub> = failure **then return** failure  
**return** [**if**  $s_1$  **then** *plan*<sub>1</sub> **else if**  $s_2$  **then** *plan*<sub>2</sub> **else** ... **if**  $s_{n-1}$  **then** *plan*<sub>*n-1*</sub>  
    **else** *plan*<sub>*n*</sub>]

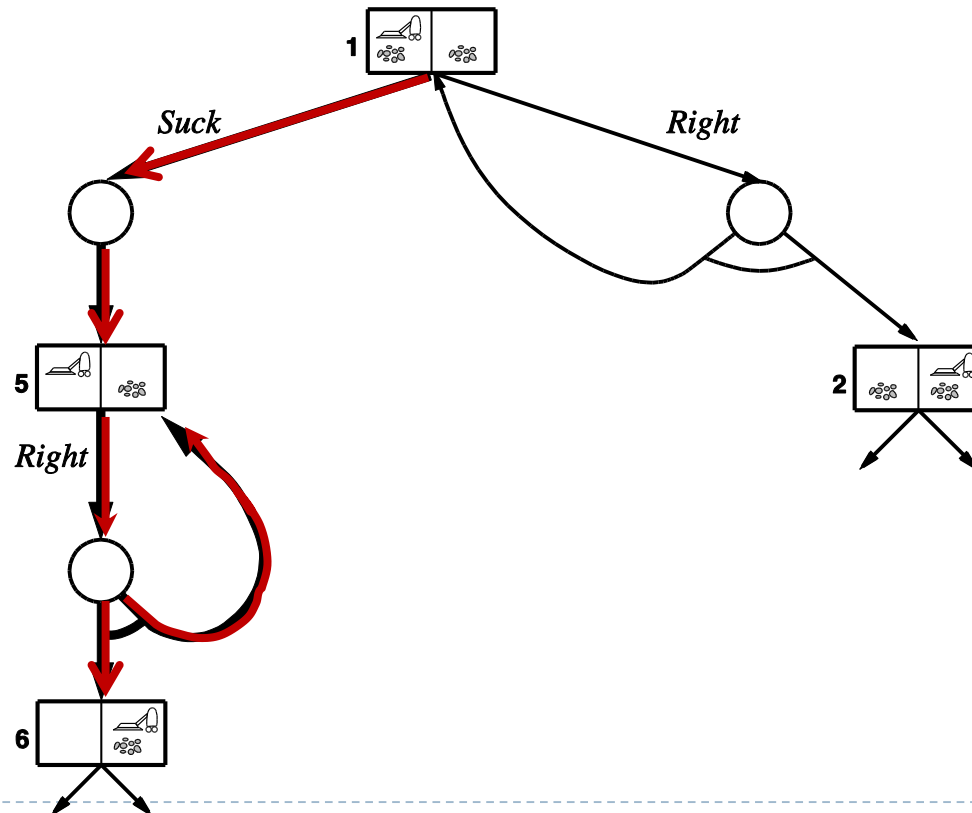
# AND-OR-GRAPH-SEARCH

---

- ▶ Cycles arise often in non-deterministic problems
  - ▶ Algorithm returns with failure when the current state is identical to one of ancestors
    - ▶ If there is a non-cyclic path, the earlier consideration of the state is sufficient
    - ▶ Termination is guaranteed in finite state spaces
      - Every path reaches a goal, a dead-end, or a repeated state


# Cycles

- ▶ Slippery vacuum world: Left and Right actions sometimes fail (leaving the agent in the same location)
- ▶ No acyclic solution



# Cycles solution

- ▶ Solution?
  - ▶ Cyclic plan: keep on trying an action until it works.
    - ▶ [Suck,  $L_1$ : Right, **if** state = 5 **then**  $L_1$  **else** Suck]
      - Or equivalently [Suck, **while** state = 5 **do** Right, Suck]
  - ▶ What changes are required in the algorithm to find cyclic solutions?



# Searching with partial observations

---

- ▶ The agent does not always know its exact state.
  - ▶ Agent is in one of several possible states and thus an action may lead to one of several possible outcomes
- ▶ **Belief state:** agent's current belief about the possible states, given the sequence of actions and observations up to that point.

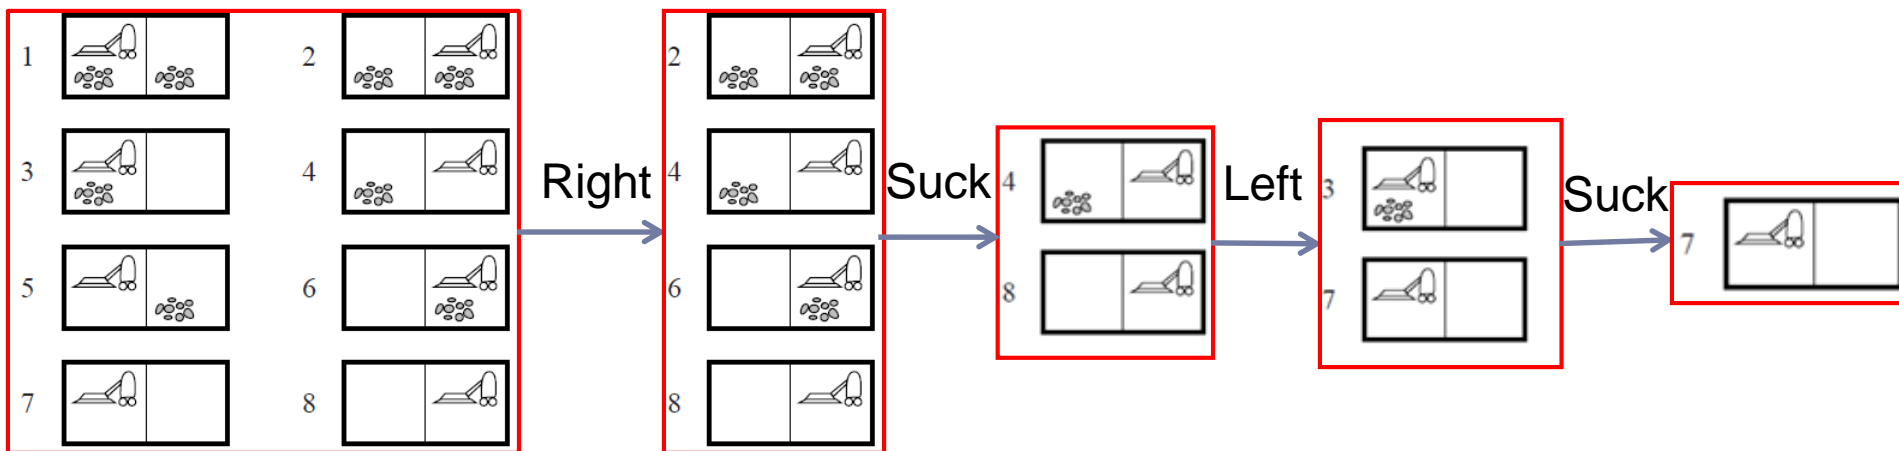
# Searching with unobservable states (Sensor-less or conformant problem)

- ▶ Initial state:

- ▶ belief = {1, 2, 3, 4, 5, 6, 7, 8}

- ▶ Action sequence (conformant plan)

- ▶ [Right, Suck, Left, Suck]



# Belief State

---

- ▶ Belief state space (instead of physical state space)
  - ▶ It is fully observable
- ▶ Physical problem:  $N$  states,  $ACTIONS_P$ ,  $RESULTS_P$ ,  $GOAL\_TEST_P$ ,  $STEP\_COST_P$
- ▶ Sensor-less problem: Up to  $2^N$  states,  $ACTIONS$ ,  $RESULTS$ ,  $GOAL\_TEST$ ,  $STEP\_COST$

# Sensor-less problem formulation (Belief-state space)


---

- ▶ **States**: every possible set of physical states,  $2^N$
- ▶ **Initial State**: usually the set of all physical states
- ▶ **Actions**:  $ACTIONS(b) = \bigcup_{s \in b} ACTIONS_P(s)$ 
  - ▶ Illegal actions?! i.e.,  $b = \{s_1, s_2\}, ACTIONS_P(s_1) \neq ACTIONS_P(s_2)$ 
    - ▶ Illegal actions have no effect on the env. (union of physical actions)
    - ▶ Illegal actions are not legal at all (intersection of physical actions)
- ▶ Solution is a sequence of actions (even in non-deterministic environment)



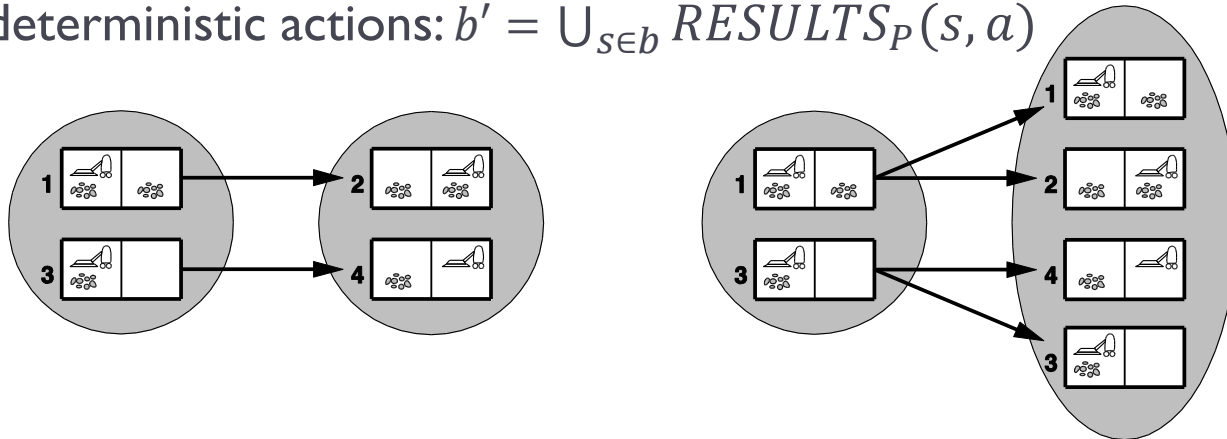
# Sensor-less problem formulation (Belief-state space)

- ▶ **Transposition model** ( $b' = PREDICT_P(b, a)$ )
  - ▶ Deterministic actions:  $b' = \{s' : s' = RESULTS_P(s, a) \text{ and } s \in b\}$
  - ▶ Nondeterministic actions:  $b' = \bigcup_{s \in b} RESULTS_P(s, a)$



# Sensor-less problem formulation (Belief-state space)

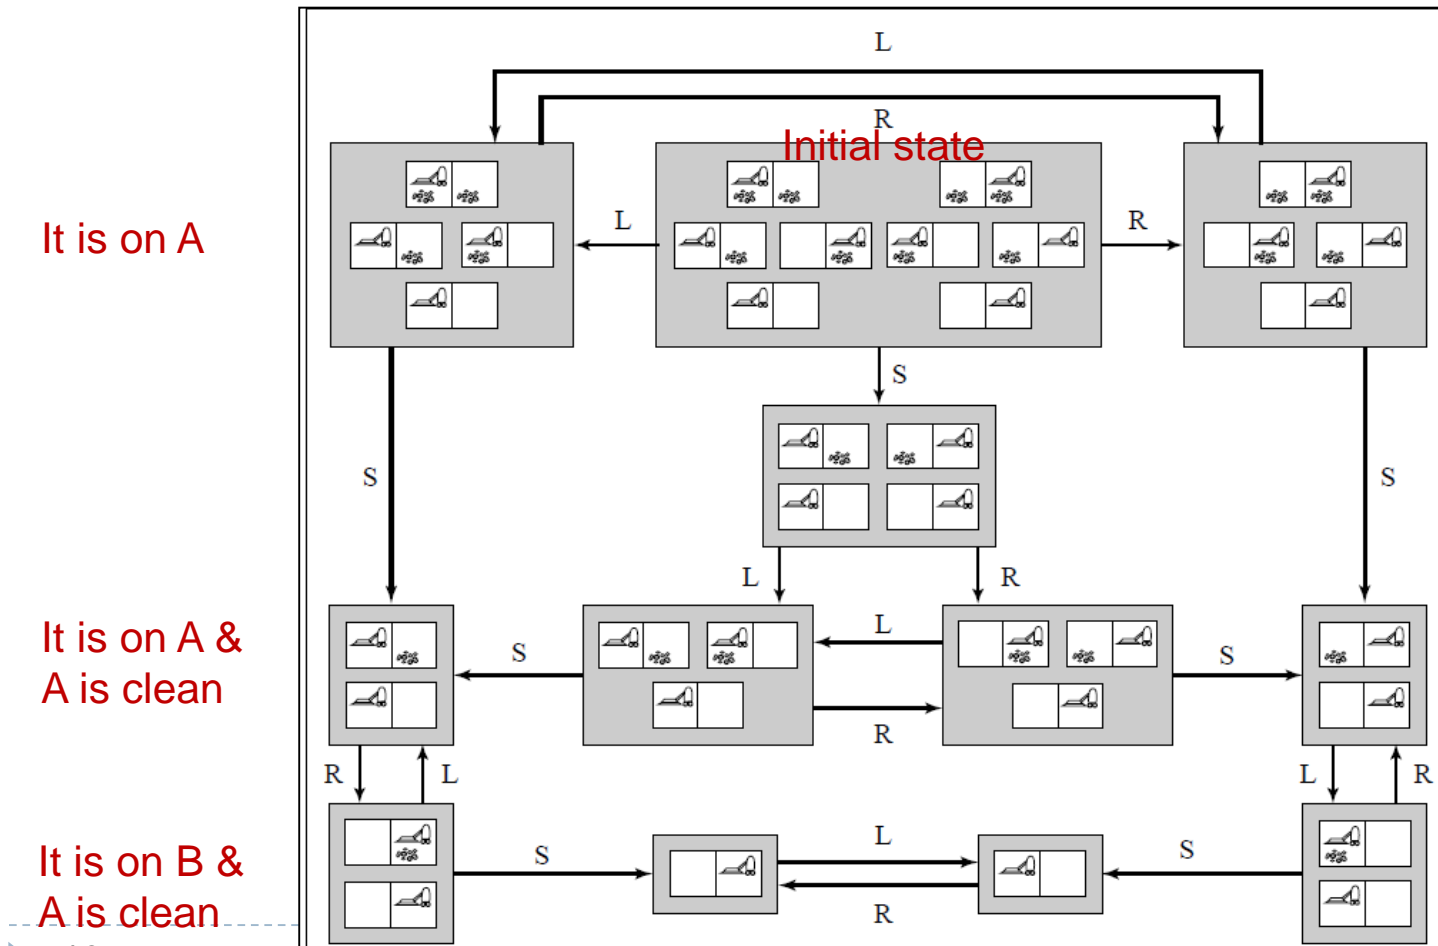
- ▶ **Transposition model** ( $b' = PREDICT_P(b, a)$ )
  - ▶ Deterministic actions:  $b' = \{s' : s' = RESULTS_P(s, a) \text{ and } s \in b\}$
  - ▶ Nondeterministic actions:  $b' = \bigcup_{s \in b} RESULTS_P(s, a)$



- ▶ **Goal test:** Goal is satisfied when all the physical states in the belief state satisfy  $GOAL\_TEST_P$ .
- ▶ **Step cost:**  $STEP\_COST_P$  if the cost of an action is the same in all states

# Belief-state space for sensor-less deterministic vacuum world

- ▶ Total number of possible belief states?  $2^8$
- ▶ Number of reachable belief states? 12



# Sensor-less problem: searching

---

- ▶ In general, we can use any standard search algorithm.
- ▶ Searching in these spaces is not usually feasible (scalability)
  - ▶ **Problem 1: No. of reachable belief states**
    - ▶ Pruning (subsets or supersets) can reduce this difficulty.
    - ▶ Branching factor and solution depth in the belief-state space and physical state space are not usually such different
  - ▶ **Problem 2 (main difficulty): No. of physical states in each belief state**
    - ▶ Using a compact state representation (like formal representation)
    - ▶ Incremental belief-state search: Search for solutions by considering physical states incrementally (not whole belief space) to quickly detect failure if we reach an unsolvable physical state.

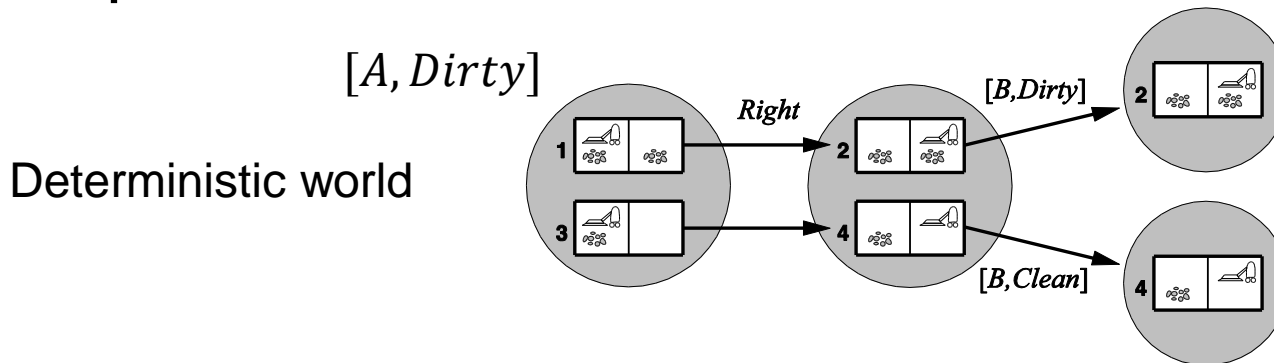
# Searching with partial observations

---

- ▶ Similar to sensor-less, after each action the new belief state must be **predicted**
- ▶ After each perception the belief state is **updated**
  - ▶ E.g., local sensing vacuum world
    - ▶ After each perception, the belief state can contain at most two physical states.
- ▶ We must plan for different **possible perceptions**

# Searching with partial observations

## A position sensor & local dirt sensor



# Transition model (partially observable env.)

---

- ▶ Prediction stage: How does the belief state change after doing an action?

$$\hat{b} = \text{PREDICT}_P(b, a)$$

- ▶ Deterministic actions:  $\hat{b} = \{s' : s' = \text{RESULTS}_P(s, a) \text{ and } s \in b\}$
- ▶ Nondeterministic actions:  $\hat{b} = \bigcup_{s \in b} \text{RESULTS}_P(s, a)$

- ▶ Possible Perceptions: What are the possible perceptions in a belief state?

$$\text{POSSIBLE\_PERCEPTS}(\hat{b}) = \{o : o = \text{PERCEPT}(s) \text{ and } s \in \hat{b}\}$$

- ▶ Update stage: How is the belief state updated after a perception?


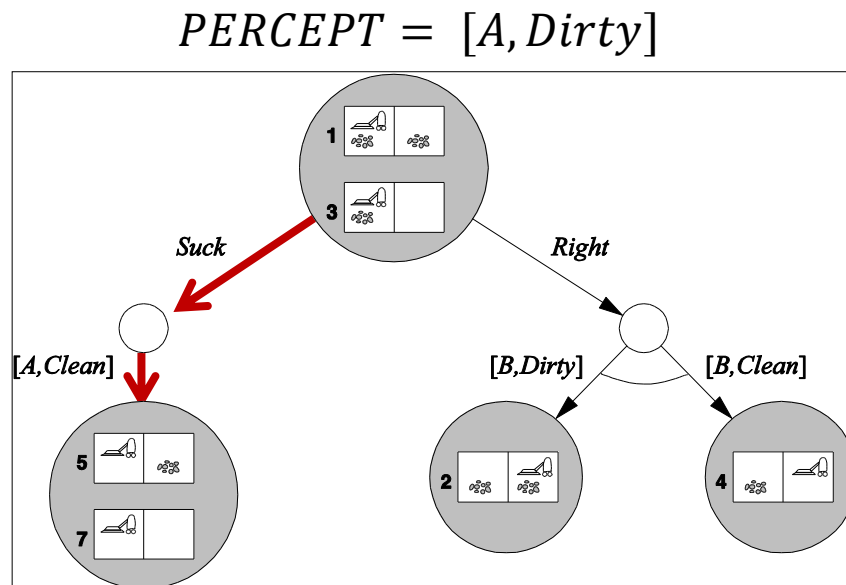
$$\text{UPDATE}(\hat{b}, o) = \{s : o = \text{PERCEPT}(s) \text{ and } s \in \hat{b}\}$$

$$\text{RESULTS}(b, a) = \{b_o : b_o = \text{UPDATE}(\text{PREDICT}(b, a), o) \text{ and } o \in \text{POSSIBLE\_PERCEPTS}(\text{PREDICT}(b, a))\}$$

# AND-OR search tree

## local sensing vacuum world

- ▶ AND-OR search tree on belief states
- ▶ First level



- ▶ Complete plan  
[Suck, Right, **if** Bstate={6} **then** Suck **else** []]



# Solving partially observable problems

---


- ▶ AND-OR graph search
- ▶ Execute the obtained contingency plan
  - ▶ Based on the achieved perception either then-part or else-part of a condition is run
  - ▶ Agent's belief state is updated when performing actions and receiving percepts
    - ▶ Maintaining the belief state is a core function of any intelligent system

$$b' = UPDATE(PREDICT(b, a), o)$$

# Kindergarten vacuum world example

## Belief state maintenance

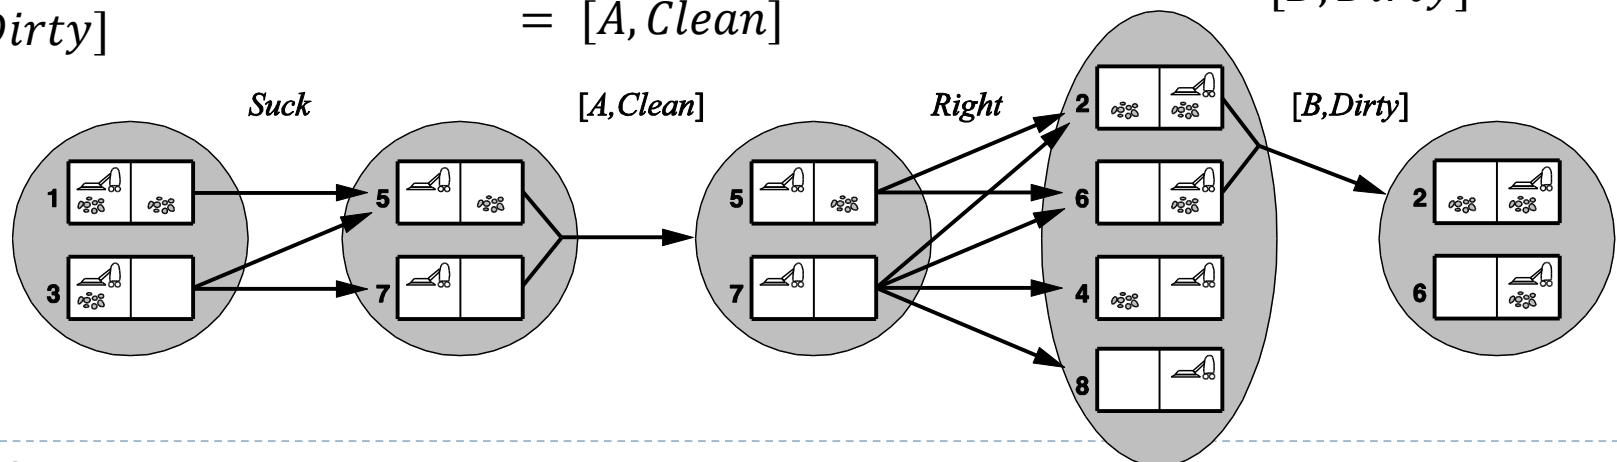
- ▶ Local sensing
- ▶ Any square may be dirty at any time (unless the agent is now cleaning it)



$$\text{PERCEPT}(s) = [A, \text{Dirty}]$$

$$\text{PERCEPT}(s) = [A, \text{Clean}]$$

$$\text{PERCEPT}(s) = [B, \text{Dirty}]$$




# Robot localization example

- ▶ **Determining current location** given a map of the world and a sequence of percepts and actions
- ▶ **Perception:** one sonar sensor in each direction (telling obstacle existence)
  - ▶ E.g., percepts=NW means there are obstacles to the north and west
- ▶ **Broken navigational system**
  - ▶ Move action randomly chooses among {Right, Left, Up, Down}


	o	o	o	o		o	o	o	o	o		o	o	o		o	
			o	o		o			o		o		o				
		o	o	o		o			o	o	o	o	o			o	
	o	o		o	o	o		o	o	o	o		o	o	o	o	

# Robot localization example (Cont.)

- ▶  $b^0$ : o squares
- ▶ Percept: NSW
- ▶  $b^1 = UPDATE(b^0, NSW)$   
(red circles)




- ▶ Execute action  $a = Move$
- ▶  $b_a^1 = PREDICT(b^1, a)$   
(red circles)



# Robot localization example (Cont.)

- ▶ Percept: NS
- ▶  $b^2 = \text{UPDATE}(b_a^1, NS)$



$\text{UPDATE}(\text{PREDICT}(\text{UPDATE}(b^0, \text{NSW}), \text{Move}), \text{NS})$

# Online search

---

- ▶ Off-line Search: solution is found before the agent starts acting in the real world
- ▶ On-line search: interleaves search and acting
  - ▶ Necessary in unknown environments
  - ▶ Useful in dynamic and semi-dynamic environments
  - ▶ Saves computational resource in non-deterministic domains (focusing only on the contingencies arising during execution)
    - ▶ Tradeoff between finding a guaranteed plan (to not get stuck in an undesirable state during execution) and required time for complete planning ahead
- ▶ Examples
  - ▶ Robot in a new environment must explore to produce a map
  - ▶ New born baby
  - ▶ Autonomous vehicles


# Online search problems

---

- ▶ Agent must perform an action to determine its outcome
  - ▶  $RESULTS(s, a)$  is found by actually being in  $s$  and doing  $a$
  - ▶ By filling  $RESULTS$  map table, the map of the environment is found.

- ▶ Different levels of ignorance

- ▶ E.g., an explorer robot may not know “laws of physics” about its actions
  - [Up, Down] action sequence gets back it to the current location



- ▶ May access to a heuristic function
- ▶ We assume deterministic & fully observable environment here
  - ▶ Also, we assume the agent knows  $ACTIONS(s)$ ,  $c(s, a, s')$  that can be used after knowing  $s'$  as the outcome,  $GOAL\_TEST(s)$

# Competitive ratio

---


- ▶ Online path cost: total cost of the path that the agent actually travels
- ▶ Best cost: cost of the shortest path “if it knew the search space in advance”
- ▶ **Competitive ratio** = Online cost / Best cost
  - ▶ Smaller values are more desirable
- ▶ Competitive ratio may be infinite
  - ▶ Dead-end state: no goal state is reachable from it
    - ▶ irreversible actions can lead to a dead-end state



# Dead-end

---

- ▶ No algorithm can avoid dead-ends in all state spaces



- ▶ Simplifying assumption: **Safely explorable** state space
  - ▶ A goal state is achievable from every reachable state

# Online search vs. offline search

---

- ▶ **Offline search:** node expansion is a simulated process rather than exerting a real action
  - ▶ Can expand a node somewhere in the state space and immediately expand a node elsewhere
- ▶ **Online search:** can discover successors only for the physical current node
  - ▶ Expand nodes in a local order
  - ▶ Interleaving search & execution

# Online search agents

---

## ▶ Online DFS

- ▶ Physical backtrack (works only for reversible actions)
  - ▶ Goes back to the state from which the agent most recently entered the current state
  - ▶ Works only for state spaces with reversible actions

## ▶ Online local search: hill-climbing

- ▶ Random walk instead of random restart
  - ▶ Randomly selecting one of available actions (preference to untried actions)
- ▶ Adding Memory (Learning Real Time A\*): more effective
  - ▶ To remember and update the costs of all visited nodes.

**function** **ONLINE-DFS**( $s'$ ) **returns** an action

inputs:  $s'$ , a percept that identifies the current state

persistent: *result*, a table indexed by state and action, initially empty

*untried*, a table that lists for each state the actions not yet tried

*unbacktracked*, a table that lists for each state the untried backtracks

$s, a$ , the previous state and action, initially null

**if** GOAL-TEST( $s'$ ) **then return** stop

**if**  $s'$  is a new state (not in *tried*) **then return**  $untried[s'] \leftarrow \text{ACTIONS}(s')$

**if**  $s$  is not null **then**

$result[s, a] \leftarrow s'$

add  $s$  to the front of *unbacktracked*[ $s'$ ]

**if** *untried*[ $s'$ ] is empty **then**

**if** *unbacktracked*[ $s'$ ] is empty **then return** stop

**else**  $a \leftarrow$  an action  $b$  such that  $result[s', b] = \text{POP}(unbacktracked[s'])$

**else**  $a \leftarrow \text{POP}(untried[s'])$

$s' \leftarrow s$

**return**  $a$