

Adversarial Search

CE417: Introduction to Artificial Intelligence
Sharif University of Technology
Spring 2016

Soleymani

“Artificial Intelligence: A Modern Approach”, 3rd Edition, Chapter 5

Outline

- ▶ Game as a search problem
- ▶ Minimax algorithm
- ▶ α - β Pruning: ignoring a portion of the search tree
- ▶ Time limit problem
 - ▶ Cut off & Evaluation function

Games as search problems

- ▶ Games
 - ▶ **Adversarial search** problems (goals are in conflict)
 - ▶ Competitive multi-agent environments
- ▶ Games in AI are a specialized kind of games (in the game theory)

Primary assumptions

- ▶ Common games in AI:
 - ▶ Two-player
 - ▶ Turn taking
 - ▶ agents act alternately
 - ▶ Zero-sum
 - ▶ agents' goals are in conflict: sum of utility values at the end of the game is zero or constant
 - ▶ Deterministic
 - ▶ Perfect information
 - ▶ fully observable




Game as a kind of search problem

- ▶ Initial state S_0 , set of states (each state contains also the turn), $ACTIONS(s)$, $RESULTS(s, a)$ like standard search
- ▶ $PLAYERS(s)$: Defines which player takes turn in a state
- ▶ $TERMINAL_TEST(s)$: Shows where game has ended
- ▶ $UTILITY(s, p)$: utility or payoff function $U: S \times P \rightarrow \mathbb{R}$ (how good is the terminal state s for player p)
 - ▶ Zero-sum (constant-sum) game: the total payoff to all players is zero (or constant) for every terminal state
 - ▶ We have utilities at end of game instead of sum of action costs


Game tree (tic-tac-toe)

- ▶ Two players: P_1 and P_2 (P_1 is now searching to find a good move)
 - ▶ Zero-sum games: P_1 gets $U(t)$, P_2 gets $C - U(t)$ for terminal node t




Game tree (tic-tac-toe)

- ▶ Two players: P_1 and P_2 (P_1 is now searching to find a good move)
 - ▶ Zero-sum games: P_1 gets $U(t)$, P_2 gets $C - U(t)$ for terminal node t



Optimal play

- ▶ Opponent is assumed optimal
- ▶ Minimax function is used to find the utility of each state.
 - ▶ MAX/MIN wants to maximize/minimize the terminal payoff




MAX gets $U(t)$ for terminal node t

Minimax

$$MINIMAX(s) = \begin{cases} UTILITY(s, MAX) & \text{if } TERMINAL_TEST(s) \\ \max_{a \in ACTIONS(s)} MINIMAX(RESULT(s, a)) & \text{PLAYER}(s) = MAX \\ \min_{a \in ACTIONS(s)} MINIMAX(RESULT(s, a)) & \text{PLAYER}(s) = MIN \end{cases}$$


Utility of being in state s

- ▶ $MINIMAX(s)$ shows the best achievable outcome of being in state s (assumption: optimal opponent)



Minimax (Cont.)

- ▶ Optimal strategy: move to the state with highest **minimax** value
 - ▶ Best achievable payoff against best play
 - ▶ Maximizes the worst-case outcome for MAX
 - ▶ It works for zero-sum games



Minimax algorithm

Depth first search

```
function MINIMAX_DECISION(state) returns an action  
  return   argmaxa ∈ ACTIONS(state) MIN_VALUE(RESULT(state, a))
```

```
function MAX_VALUE(state) returns a utility value  
  if TERMINAL_TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN\_VALUE}(\text{RESULTS}(\text{state}, a)))$   
  return  $v$ 
```

```
function MIN_VALUE(state) returns a utility value  
  if TERMINAL_TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX\_VALUE}(\text{RESULTS}(\text{state}, a)))$   
  return  $v$ 
```


Properties of minimax

- ▶ Complete? Yes (when tree is finite)
- ▶ Optimal? Yes (against an optimal opponent)
- ▶ Time complexity: $O(b^m)$
- ▶ Space complexity: $O(bm)$ (depth-first exploration)
- ▶ For chess, $b \approx 35$, $m > 50$ for reasonable games
 - ▶ Finding exact solution is completely infeasible


Pruning

- ▶ Correct minimax decision without looking at every node in the game tree
 - ▶ α - β pruning
 - ▶ Branch & bound algorithm
 - ▶ Prunes away branches that cannot influence the final decision


α - β pruning example




α - β pruning example




α - β pruning example









α - β pruning example



α - β pruning example




α - β progress



α - β pruning


- ▶ Assuming depth-first generation of tree
 - ▶ We prune node n when player has a better choice m at (parent or) any ancestor of n
- ▶ Two types of pruning (cuts):
 - ▶ pruning of max nodes (α -cuts)
 - ▶ pruning of min nodes (β -cuts)



Why is it called α - β ?

- ▶ α : Value of the best (highest) choice found so far at any choice point along the path for **MAX**
- ▶ β : Value of the best (lowest) choice found so far at any choice point along the path for **MIN**
- ▶ Updating α and β during the search process
- ▶ For a MAX node once the value of this node is known to be more than the current β ($v \geq \beta$), its remaining branches are pruned.
- ▶ For a MIN node once the value of this node is known to be less than the current α ($v \leq \alpha$), its remaining branches are pruned.

α - β pruning (an other example)




function *ALPHA_BETA_SEARCH*(state) **returns** an action
 $v \leftarrow \text{MAX_VALUE}(\text{state}, -\infty, +\infty)$
 return the action **in** *ACTIONS*(state) with value v

function *MAX_VALUE*(state, α , β) **returns** a utility value
 if *TERMINAL_TEST*(state) **then return** *UTILITY*(state)
 $v \leftarrow -\infty$
 for each a **in** *ACTIONS*(state) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN_VALUE}(\text{RESULTS}(\text{state}, a), \alpha, \beta))$
 if $v \geq \beta$ **then return** v
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
 return v

function *MIN_VALUE*(state, α , β) **returns** a utility value
 if *TERMINAL_TEST*(state) **then return** *UTILITY*(state)
 $v \leftarrow +\infty$
 for each a **in** *ACTIONS*(state) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX_VALUE}(\text{RESULTS}(\text{state}, a), \alpha, \beta))$
 if $v \leq \alpha$ **then return** v
 $\beta \leftarrow \text{MIN}(\beta, v)$
 return v

Order of moves

- ▶ Good move ordering improves effectiveness of pruning



- ▶ Best order: time complexity is $O(b^{m/2})$
- ▶ Random order: time complexity is about $O(b^{3m/4})$ for moderate b
 - ▶ α - β pruning just improves the search time only partly

Computational time limit (example)

- ▶ 100 secs is allowed for each move (game rule)
- ▶ 10^4 nodes/sec (processor speed)
- ▶ We can explore just 10^6 nodes for each move
 - ▶ $b^m = 10^6, b=35 \Rightarrow m=4$
(4-ply look-ahead is a hopeless chess player!)

Computational time limit: Solution

- ▶ We must make a decision even when finding the optimal move is infeasible.
- ▶ Cut off the search and apply a heuristic evaluation function
 - ▶ **cutoff test:** turns non-terminal nodes into terminal leaves
 - ▶ Cut off test instead of terminal test (e.g., depth limit)
 - ▶ **evaluation function:** estimated desirability of a state
 - ▶ Heuristic function evaluation instead of utility function
- ▶ This approach does not guarantee optimality.

Heuristic minimax

$$H_{MINIMAX}(s,d) =$$

$$\begin{cases} EVAL(s, MAX) & \text{if } CUTOFF_TEST(s, d) \\ \max_{a \in ACTIONS(s)} H_{MINIMAX}(RESULT(s, a), d + 1) & \text{PLAYER}(s) = MAX \\ \min_{a \in ACTIONS(s)} H_{MINIMAX}(RESULT(s, a), d + 1) & \text{PLAYER}(s) = MIN \end{cases}$$

Evaluation functions

- ▶ For terminal states, it should order them in the same way as the true utility function.
- ▶ For non-terminal states, it should be strongly correlated with the actual **chances of winning**.
- ▶ It must not need high computational cost.

Evaluation functions based on features

- ▶ **Example: features for evaluation of the chess states**
 - ▶ Number of each kind of piece: number of white pawns, black pawns, white queens, black queens, etc
 - ▶ King safety
 - ▶ Good pawn structure

Evaluation functions

- ▶ Weighted sum of features

- ▶ Assumption: contribution of each feature is independent of the value of the other features

$$EVAL(s) = w_1 \times f_1(s) + w_2 \times f_2(s) + \cdots + w_n \times f_n(s)$$

- ▶ Weights can be assigned based on the human experience or machine learning methods.

- ▶ Example: Chess

- ▶ Features: number of white pawns (f_1), number of white bishops (f_2), number of white rooks (f_3), number of black pawns (f_4), ...
 - ▶ Weights: $w_1 = 1, w_2 = 3, w_3 = 5, w_4 = -1, \dots$

Cutting off search: simple depth limit

- ▶ Simple: depth limit d_0

$$CUTOFF_TEST(s, d) = \begin{cases} true & \text{if } d > d_0 \text{ or } TERMINAL_TEST(s) = TRUE \\ false & \text{otherwise} \end{cases}$$


Cutting off search: simple depth limit

▶ Problem 1: **non-quiet** positions

- ▶ Few more plies make big difference in evaluation value

▶ Problem 2: **horizon effect**

- ▶ Delaying tactics against opponent's move that causes serious unavoidable damage (because of pushing the damage beyond the horizon that the player can see)



More sophisticated cutting off

- ▶ Cutoff only on quiescent positions
 - ▶ **Quiescent search**: expanding non-quiescent positions until reaching quiescent ones
- ▶ Horizon effect
 - ▶ **Singular extension**: a move that is clearly better than all other moves in a given position.
 - ▶ Once reaching the depth limit, check to see if the singular extension is a legal move.
 - ▶ It makes the tree deeper but it does not add many nodes to the tree due to few possible singular extensions.

Speed up the search process

- ▶ **Table lookup** rather than search for some states
 - ▶ E.g., for the opening and ending of games (where there are few choices)
- ▶ **Example: Chess**
 - ▶ For each opening, the best advice of human experts (from books describing good plays) can be copied into tables.
 - ▶ For endgame, computer analysis is usually used (solving endgames by computer).

Stochastic games: Backgammon

MAX


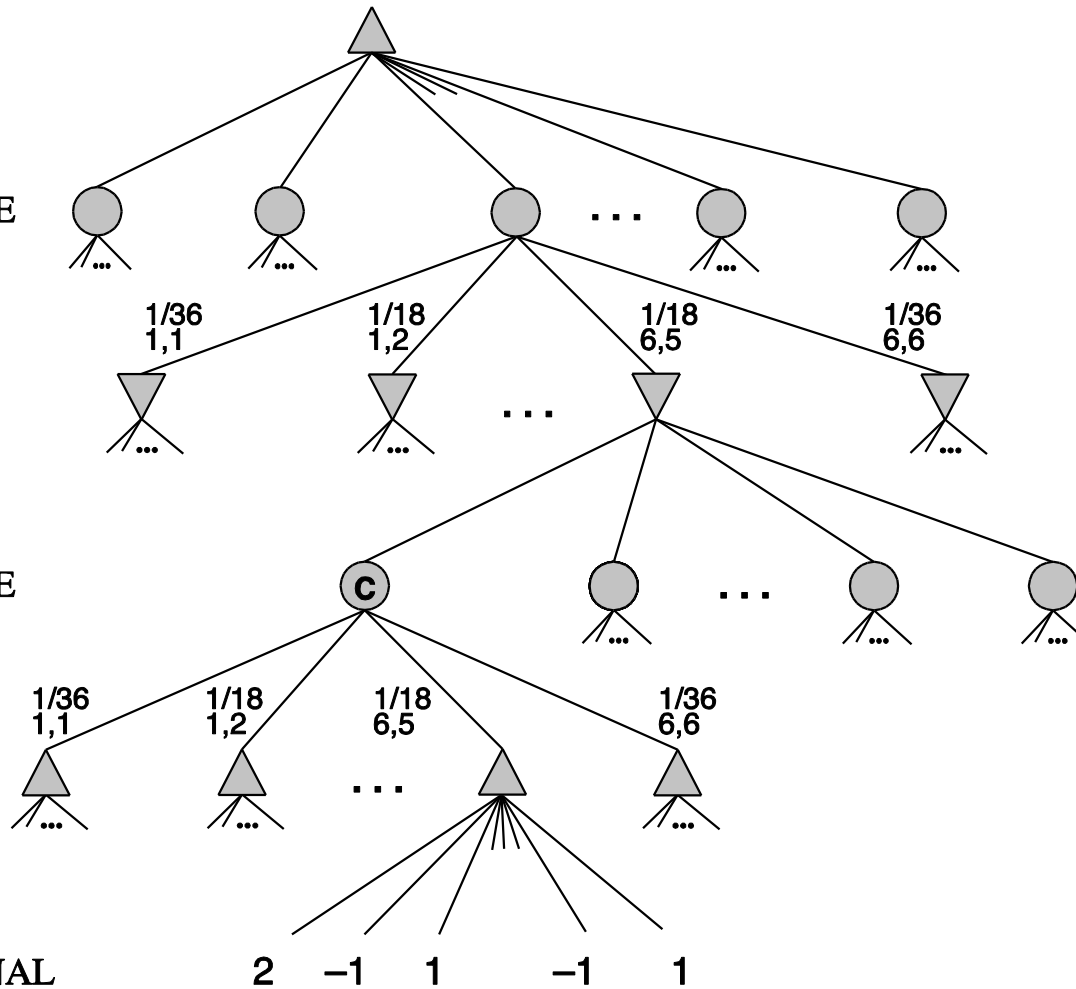
CHANCE

MIN

CHANCE


MAX

TERMINAL



Stochastic games

- ▶ **Expected utility:** Chance nodes take average (expectation) over all possible outcomes.
- ▶ It is consistent with the definition of rational agents trying to maximize expected utility.




Stochastic games

$$EXPECT_MINIMAX(s) = \begin{cases} UTILITY(s, MAX) & \text{if } TERMINAL_TEST(s) \\ \max_{a \in ACTIONS(s)} EXPECT_MINIMAX(RESULT(s, a)) & PLAYER(s) = MAX \\ \min_{a \in ACTIONS(s)} EXPECT_MINIMAX(RESULT(s, a)) & PLAYER(s) = MIN \\ \sum_r P(r) EXPECT_MINIMAX(RESULT(s, r)) & PLAYER(s) = CHANCE \end{cases}$$

Evaluation functions for stochastic games

- ▶ An order preserving transformation on leaf values is not a sufficient condition.



- ▶ Evaluation function must be a positive linear transformation of the expected utility of a position.

Properties of search space for stochastic games

- ▶ $O(b^m n^m)$
 - ▶ Backgammon: $b \approx 20$ (can be up to 4000 for double dice rolls), $n = 21$ (no. of different possible dice rolls)
 - ▶ 3-ply is manageable ($\approx 10^8$ nodes)
- ▶ Probability of reaching a given node decreases enormously by increasing the depth (multiplying probabilities)
 - ▶ Forming detailed plans of actions may be pointless
 - ▶ Limiting depth is not such damaging particularly when the probability values (for each non-deterministic situation) are close to each other
 - ▶ But pruning is not straightforward.

Search algorithms for stochastic games

- ▶ Advanced alpha-beta pruning

- ▶ Pruning MIN and MAX nodes as alpha-beta
- ▶ Pruning chance nodes (by putting bounds on the utility values and so placing an upper bound on the value of a chance node)


- ▶ Monte Carlo simulation to evaluate a position

- ▶ Starting from the corresponding position, the algorithm plays thousands of games against itself using random dice rolls.
 - ▶ **Win percentage** as the approximated value of the position (Backgammon)

State-of-the-art game programs


▶ Chess ($b \approx 35$)

- ▶ In 1997, Deep Blue defeated Kasparov.
 - ▶ ran on a parallel computer doing alpha-beta search.
 - ▶ reaches depth 14 plies routinely.
 - techniques to extend the effective search depth
- ▶ Hydra: Reaches depth 18 plies using more heuristics.





▶ Checkers ($b < 10$)

- ▶ Chinook (ran on a regular PC and uses alpha-beta search) ended 40-year-reign of human world champion Tinsley in 1994.
- ▶ Since 2007, Chinook has been able to **play perfectly** by using alpha-beta search combined with a database of 39 trillion endgame positions.



State-of-the-art game programs (Cont.)


- ▶ Othello (b is usually between 5 to 15)
 - ▶ Logistello defeated the human world champion by six games to none in 1997.
 - ▶ Human champions are no match for computers at Othello.
- ▶ Go ($b > 300$)
 - ▶ Human champions refuse to compete against computers (current programs are still at advanced amateur level).
 - ▶ MOGO avoided alpha-beta search and used Monte Carlo rollouts.



State-of-the-art game programs (Cont.)


▶ Backgammon (stochastic)

- ▶ TD-Gammon (1992) was competitive with top human players.
 - ▶ Depth 2 or 3 search along with a good evaluation function developed by learning methods



▶ Bridge (partially observable, multiplayer)

- ▶ In 1998, GIB was 12th in a field of 35 in the par contest at human world championship.
- ▶ In 2005, Jack defeated three out of seven top champions pairs. Overall, it lost by a small margin.



▶ Scrabble (partially observable & stochastic)

- ▶ In 2006, Quackle defeated the former world champion 3-2.

