



# Mutual Exclusion

Distributed Systems

---

**Ali Kamandi, PH.D.**

School of Engineering Science

College of Engineering

University of Tehran


**kamandi@ut.ac.ir**

**2024**



# Mutual Exclusion

- **Trying** : trying to get into critical section.
- **Termination**: in critical section.
- **Exiting** : cleaning up so that other processes can enter their critical sections.
- **Remainder** : everything else—essentially just going about its non-critical business.



# Goals

**Mutual exclusion** At most one process is in the critical state at a time.

**No deadlock (progress)** If there is at least one process in a trying state, then eventually some process enters a critical state; similarly for exiting and remainder states.

**No lockout (lockout-freedom):** If there is a particular process in a trying or exiting state, that process eventually leaves that state. This means that I don't starve because somebody else keeps jumping past me and seizing the critical resource before I can.

Stronger versions of lockout-freedom include explicit time bounds (how many rounds can go by before I get in) or **bounded bypass** (nobody gets in more than  $k$  times before I do).

# approach

- Token-based approach.
- Non-token-based approach.
- Quorum-based approach.

# Mutual exclusion using strong primitives

A **test-and-set** operation does the following sequence of actions atomically:

```
1 oldValue ← read(bit)
2 write(bit, 1)
3 return oldValue
```

```
1 while true do
  // trying
2   while TAS(lock) = 1 do nothing
  // critical
3   (do critical section stuff)
  // exiting
4   reset(lock)
  // remainder
5   (do remainder stuff)
```

# A lockout-free algorithm using an atomic queue

```
1 while true do
  // trying
2   enq(q, myld)
3   while peek(q)  $\neq$  myld do nothing
  // critical
4   (do critical section stuff)
  // exiting
5   deq(q)
  // remainder
6   (do remainder stuff)
```

# Mutual exclusion using only atomic registers

```
shared data:
1 waiting, initially arbitrary
2 present[i] for  $i \in \{0, 1\}$ , initially 0
3 Code for process  $i$ :
4 while true do
    // trying
5     present[i]  $\leftarrow$  1
6     waiting  $\leftarrow i$ 
7     while true do
8         if present[ $\neg i$ ] = 0 then
9             break
10        if waiting  $\neq i$  then
11            break
    // critical
12    (do critical section stuff)
    // exiting
13    present[i] = 0
    // remainder
14    (do remainder stuff)
```

**Peterson's** mutual exclusion algorithm for two processes

# Requirements of mutual exclusion algorithms

- **Safety property** The safety property states that at any instant, only one process can execute the critical section. This is an essential property of a mutual exclusion algorithm.
- **Liveness property** This property states the absence of deadlock and starvation. Two or more sites should not endlessly wait for messages that will never arrive. In addition, a site must not wait indefinitely to execute the CS while other sites are repeatedly executing the CS. That is, every requesting site should get an opportunity to execute the CS in finite time.
- **Fairness** Fairness in the context of mutual exclusion means that each process gets a fair chance to execute the CS. In mutual exclusion algorithms, the fairness property generally means that the CS execution requests are executed in order of their arrival in the system (the time is determined by a logical clock).

The first property is absolutely necessary and the other two properties are considered important in mutual exclusion algorithms.



# Performance Metrics

- **Message complexity** This is the number of messages that are required per CS execution by a site.
- **Synchronization delay** After a site leaves the CS, it is the time required and before the next site enters the CS). Note that normally one or more sequential message exchanges may be required after a site exits the CS and before the next site can enter the CS.

**Response time** This is the time interval a request waits for its CS execution to be over after its request messages have been sent out. Thus, response time does not include the

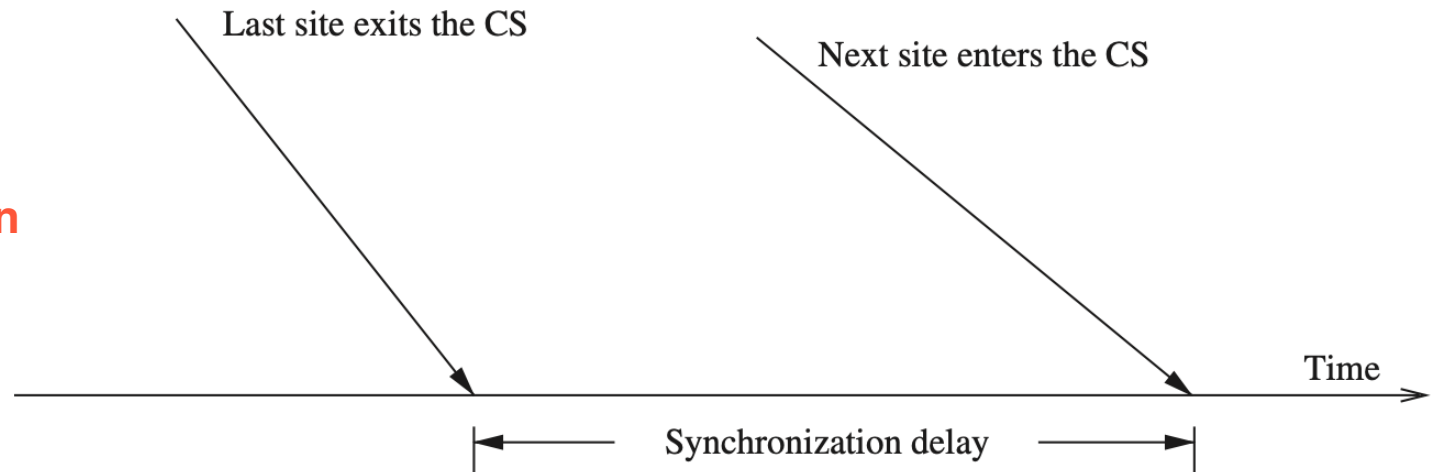
- time a request waits at a site before its request messages have been sent out. The first property is absolutely necessary and the other two properties are considered important in mutual exclusion algorithms.

**System throughput** This is the rate at which the system executes requests for the CS. If SD is the synchronization delay and E is the average critical section execution time, then the throughput is given by the following equation:

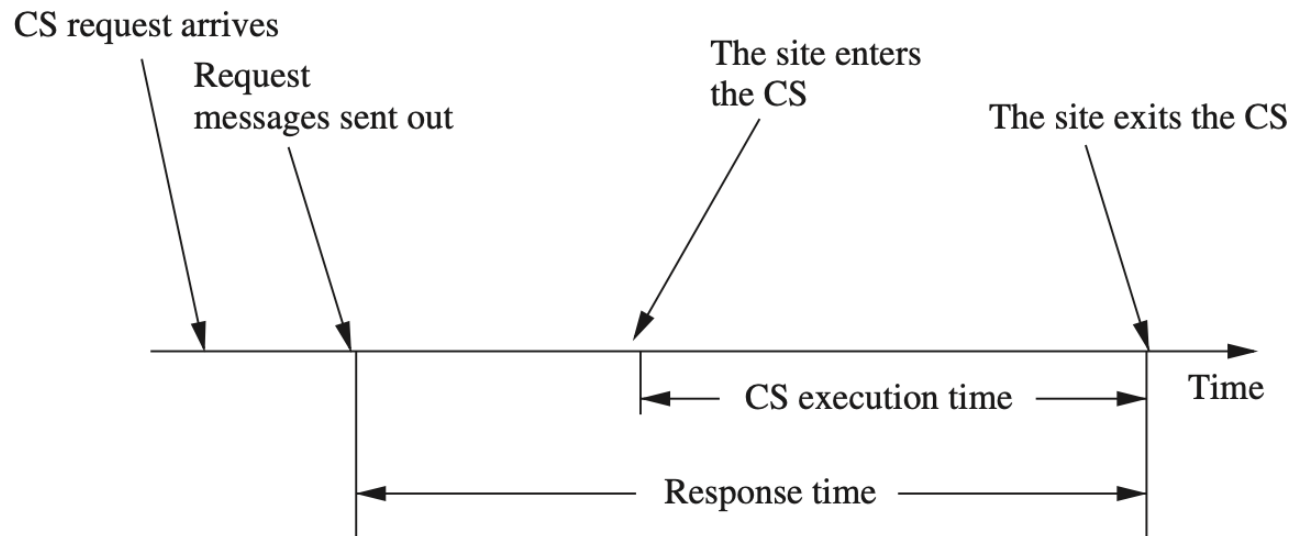
$$\text{System Throughput} = \frac{1}{(SD + E)}$$

# Performance Metrics

## Synchronization delay



## Response time



# Lamport's Algorithm

## Requesting the critical section

- When a site  $S_i$  wants to enter the CS, it broadcasts a REQUEST( $ts_i, i$ ) message to all other sites and places the request on *request\_queue<sub>i</sub>*. ( $(ts_i, i)$  denotes the timestamp of the request.)
- When a site  $S_j$  receives the REQUEST( $ts_i, i$ ) message from site  $S_i$ , it places site  $S_i$ 's request on *request\_queue<sub>j</sub>* and returns a timestamped REPLY message to  $S_i$ .

## Executing the critical section

Site  $S_i$  enters the CS when the following two conditions hold:

**L1:**  $S_i$  has received a message with timestamp larger than  $(ts_i, i)$  from all other sites.


**L2:**  $S_i$ 's request is at the top of *request\_queue<sub>i</sub>*.

## Releasing the critical section


- Site  $S_i$ , upon exiting the CS, removes its request from the top of its request queue and broadcasts a timestamped RELEASE message to all other sites.
- When a site  $S_j$  receives a RELEASE message from site  $S_i$ , it removes  $S_i$ 's request from its request queue.

# Lamport's Algorithm

Site  $S_1$  enters the **CS**.




Site  $S_1$  exits the **CS** and sends **RELEASE** messages.




# Lamport's Algorithm

Site  $S_1$  exits the **CS** and sends **RELEASE** messages.



Site  $S_2$  enters the **CS**.



# Performance

For each CS execution, Lamport's algorithm requires  **$N-1$  REQUEST messages**,  **$N-1$  REPLY messages**, and  **$N-1$  RELEASE messages**. Thus, Lamport's algorithm requires  **$3N-1$  messages per CS invocation**. The synchronization delay in the algorithm is  $T$ .

بهینه سازی: **عدم ارسال پاسخ**

کاهش تعداد پیام بین  $2(N-1)$  و  $3(N-1)$

# Ricart–Agrawala algorithm

## Requesting the critical section

- (a) When a site  $S_i$  wants to enter the CS, it broadcasts a timestamped REQUEST message to all other sites.
- (b) When site  $S_j$  receives a REQUEST message from site  $S_i$ , it sends a REPLY message to site  $S_i$  if site  $S_j$  is neither requesting nor executing the CS, or if the site  $S_j$  is requesting and  $S_i$ 's request's timestamp is smaller than site  $S_j$ 's own request's timestamp. Otherwise, the reply is deferred and  $S_j$  sets  $RD_j[i] := 1$ .


## Executing the critical section

- (c) Site  $S_i$  enters the CS after it has received a REPLY message from every site it sent a REQUEST message to.

## Releasing the critical section

- (d) When site  $S_i$  exits the CS, it sends all the deferred REPLY messages:  $\forall j$  if  $RD_i[j] = 1$ , then sends a REPLY message to  $S_j$  and sets  $RD_i[j] := 0$ .


# Ricart–Agrawala algorithm: example




Sites  $S_1$  and  $S_2$  each make a request for the **CS**.



# Ricart–Agrawala algorithm: example




# Ricart–Agrawala algorithm: example



Site  $S_1$  exits the **CS** and sends a REPLY message to  $S_2$ 's deferred request.

# Ricart–Agrawala algorithm: example



# Performance

For each CS execution, the Ricart–Agrawala algorithm requires  **$N-1$  REQUEST messages** and  **$N-1$  REPLY messages**. Thus, it requires  **$2N-1$  messages per CS execution**. The synchronization delay in the algorithm is  $T$ .


# Quorum-based mutual exclusion algorithms

## Maekawa's algorithm

- M1  $(\forall i \forall j : i \neq j, 1 \leq i, j \leq N :: R_i \cap R_j \neq \phi).$
- M2  $(\forall i : 1 \leq i \leq N :: S_i \in R_i).$
- M3  $(\forall i : 1 \leq i \leq N :: |R_i| = K).$
- M4 Any site  $S_j$  is contained in  $K$  number of  $R_i$ s,  $1 \leq i, j \leq N$ .

Maekawa used the theory of projective planes and showed that  $N=K(K-1) +1$ . This relation gives  $|R_i| = \sqrt{N}$ .

# Meakawa's algorithm



$$S_1 = \{1,2,3,4\}$$

$$S_2 = \{2,5,8,11\}$$

$$S_3 = \{3,5,9,13\}$$

$$S_4 = \{4,5,10,12\}$$

$$S_5 = \{5,1,6,7\}$$

$$S_6 = \{6,2,9,12\}$$

$$S_7 = \{7,3,8,12\}$$

$$S_8 = \{8,1,9,10\}$$

$$S_9 = \{9,4,7,11\}$$

$$S_{10} = \{10,2,7,13\}$$

$$S_{11} = \{11,3,6,10\}$$

$$S_{12} = \{12,1,11,13\}$$

$$S_{13} = \{13,4,6,8\}$$

# Meakawa's algorithm

## **Requesting the critical section:**

- (a) A site  $S_i$  requests access to the CS by sending REQUEST( $i$ ) messages to all sites in its request set  $R_i$ .
- (b) When a site  $S_j$  receives the REQUEST( $i$ ) message, it sends a REPLY( $j$ ) message to  $S_i$  provided it hasn't sent a REPLY message to a site since its receipt of the last RELEASE message. Otherwise, it queues up the REQUEST( $i$ ) for later consideration.

## **Executing the critical section:**

- (c) Site  $S_i$  executes the CS only after it has received a REPLY message from every site in  $R_i$ .

## **Releasing the critical section:**

- (d) After the execution of the CS is over, site  $S_i$  sends a RELEASE( $i$ ) message to every site in  $R_i$ .
- (e) When a site  $S_j$  receives a RELEASE( $i$ ) message from site  $S_i$ , it sends a REPLY message to the next site waiting in the queue and deletes that entry from the queue. If the queue is empty, then the site updates its state to reflect that it has not sent out any REPLY message since the receipt of the last RELEASE message.


# Performance

Note that the size of a request set is  $\sqrt{N}$ . Therefore, an execution of the CS requires  $\sqrt{N}$  **REQUEST**,  $\sqrt{N}$  **REPLY** and  $\sqrt{N}$  **RELEASE messages**, resulting in  **$3\sqrt{N}$  messages per CS execution**. Synchronization delay in this algorithm is  $2T$ . This is because after a site  $S_i$  exits the CS, it first releases all the sites in  $R_i$  and then one of those sites sends a **REPLY** message to the next site that executes the CS. Thus, two sequential message transfers are required between two successive CS executions. Maekawa's algorithm is deadlock-prone. Measures to handle deadlocks require additional messages.



# Token-based algorithms

## Raymond's tree-based algorithm



$\text{HOLDER}_B = C$

$\text{HOLDER}_C = G$

$\text{HOLDER}_D = C$

$\text{HOLDER}_E = A$

$\text{HOLDER}_F = B$

$\text{HOLDER}_G = \text{self}$

# References

- **Aspnes, Chapter 18**
- **Singhal, Chapter 9**